TEST DATA GENERATION

End-to-end scenario

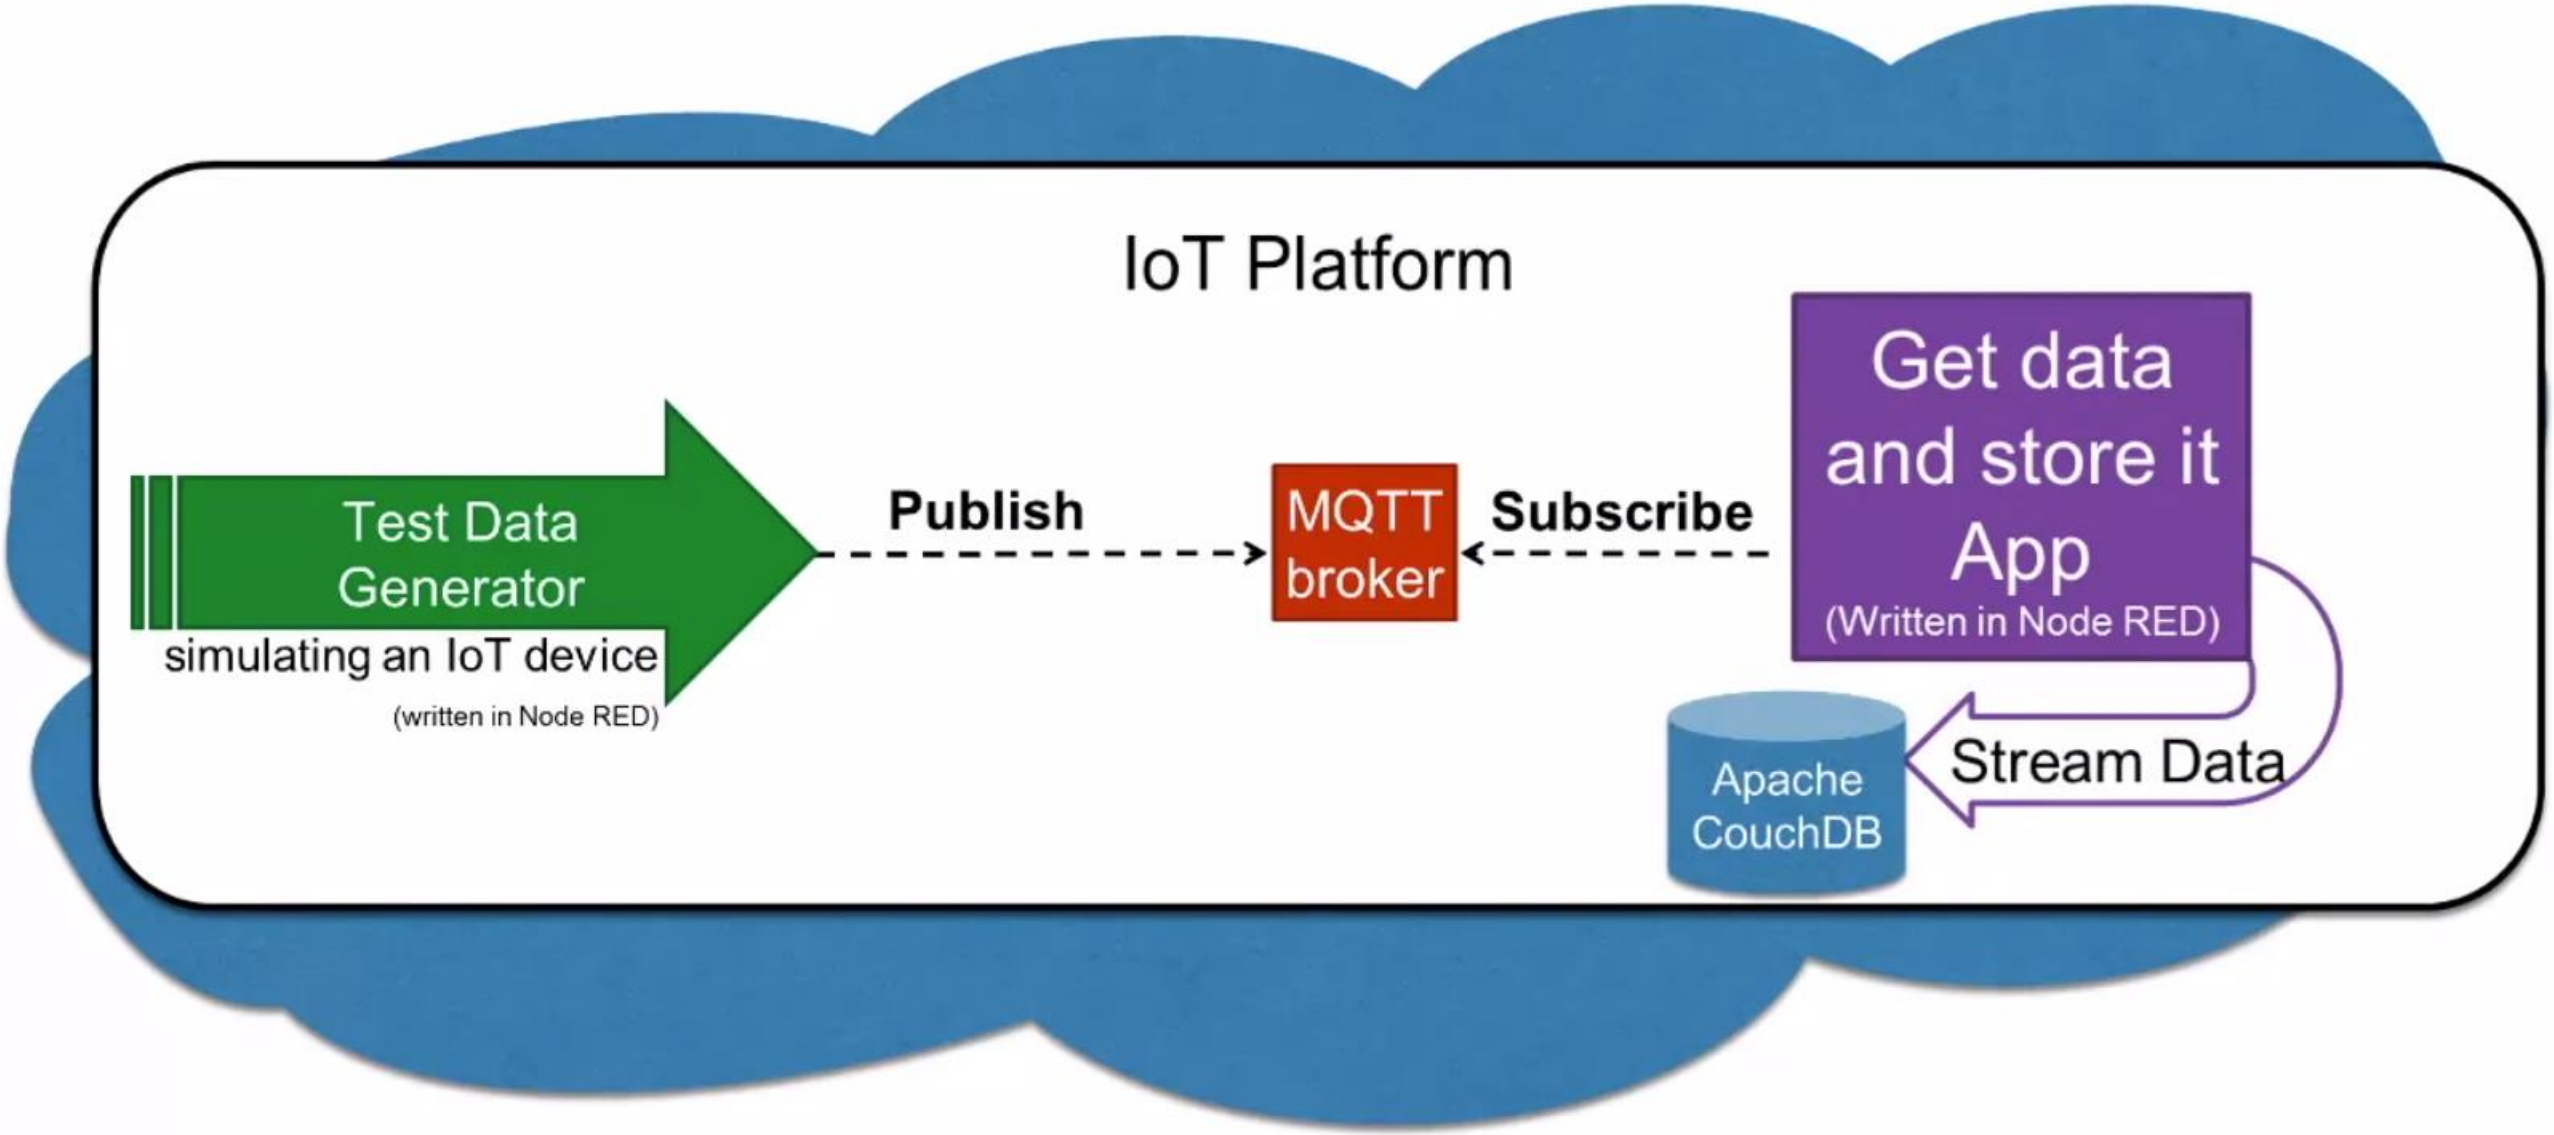

FLOW 1

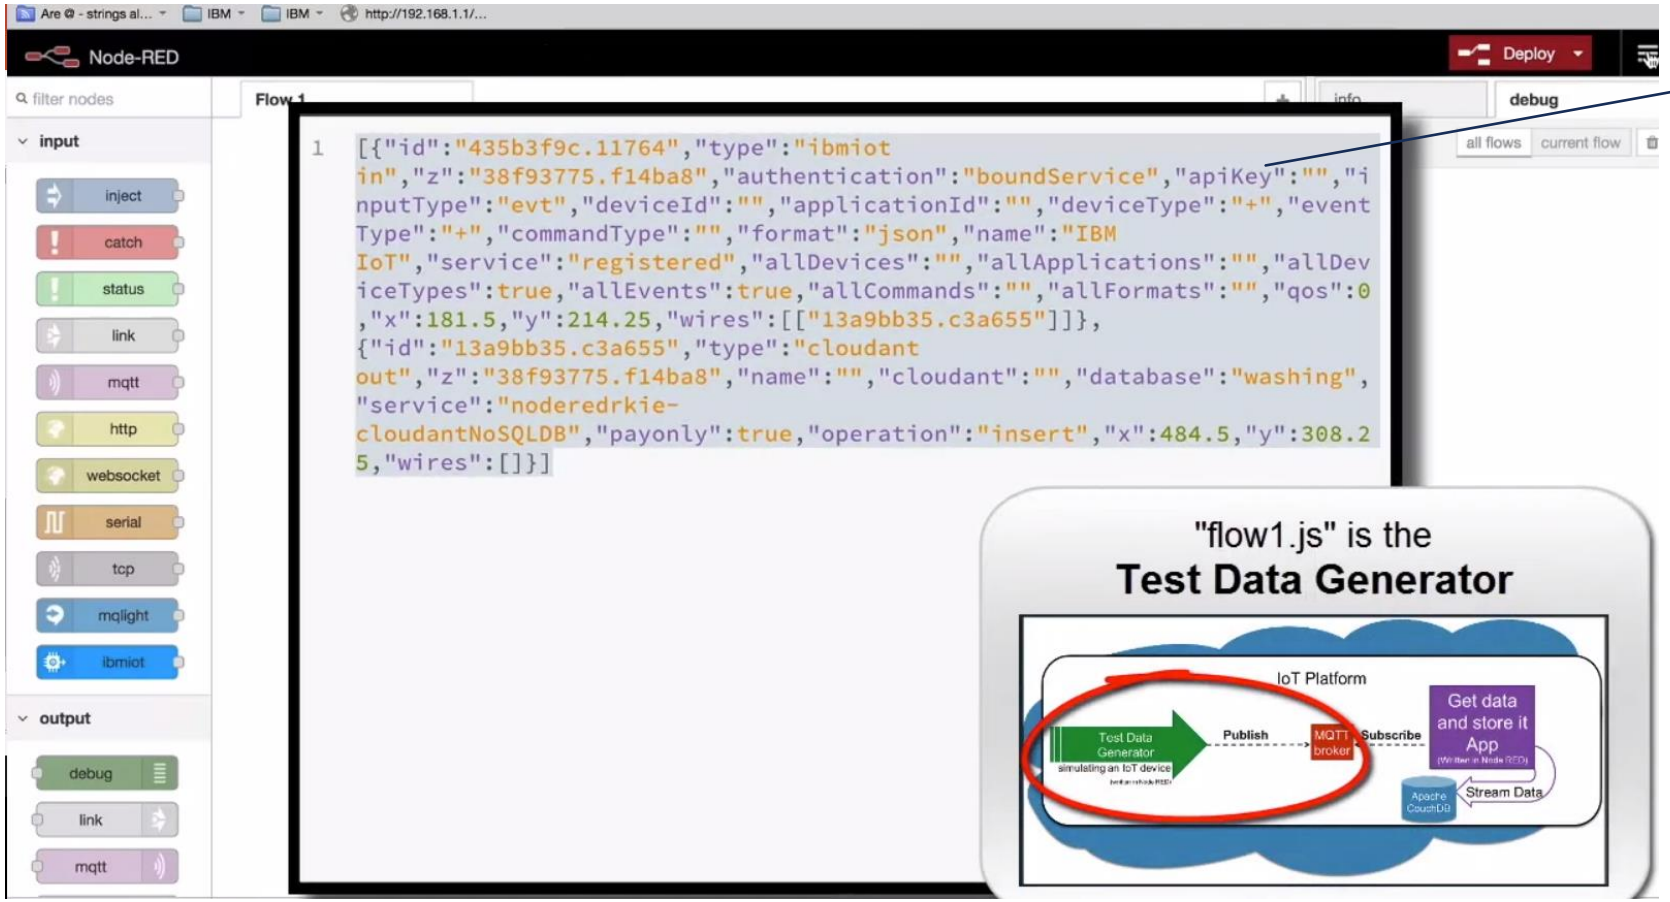

Open in a text editor

Past here:

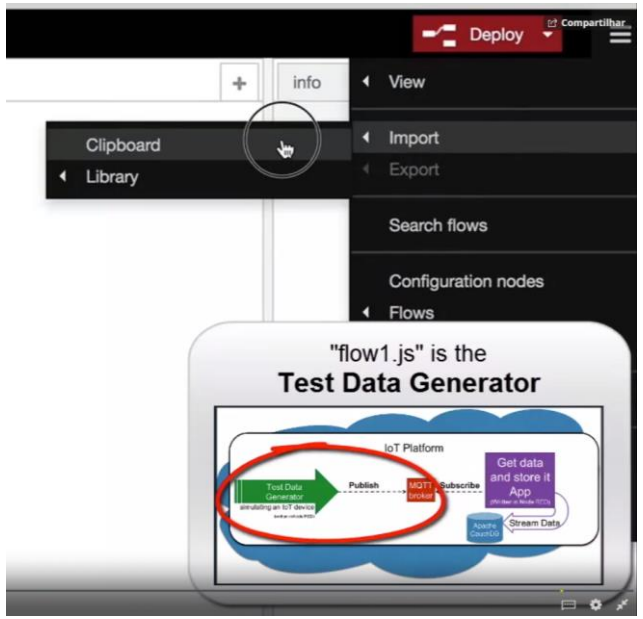

FLOW 2

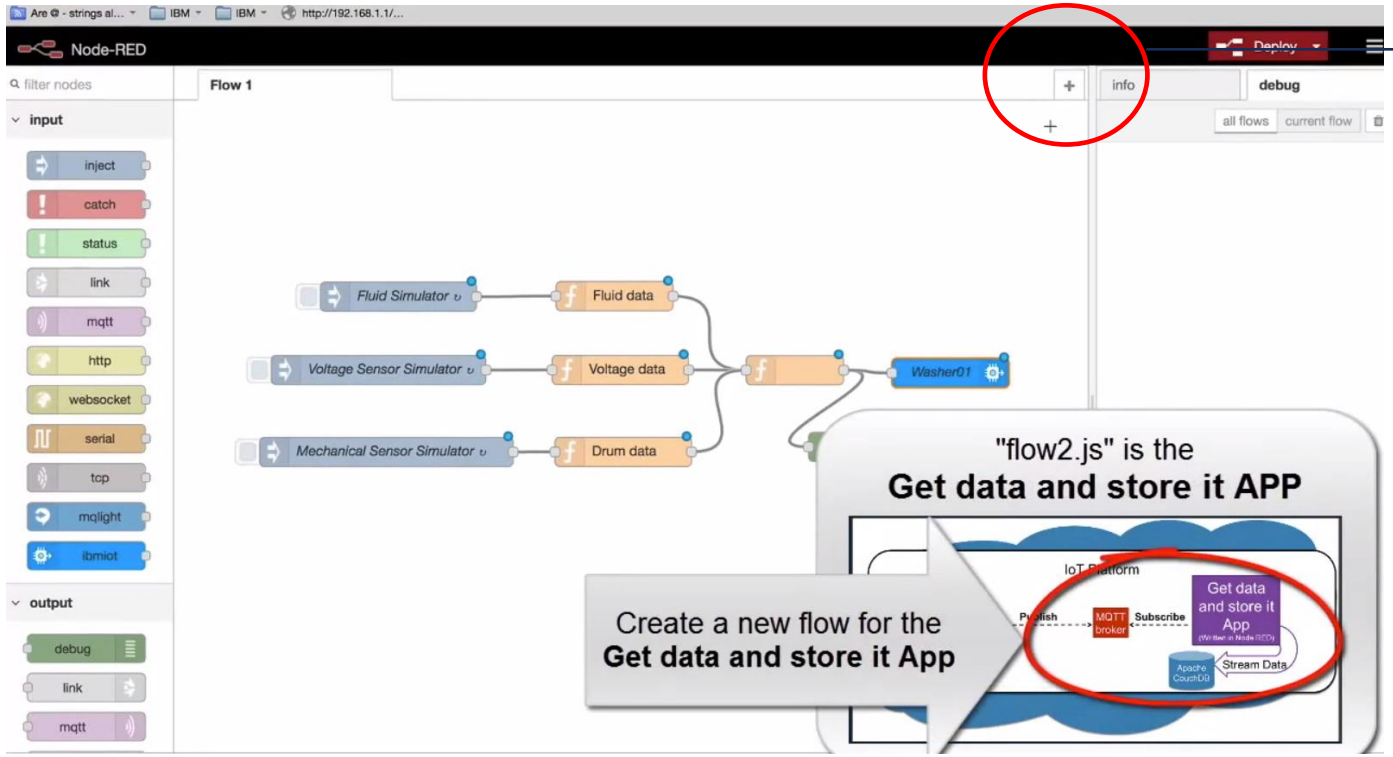

Click here to create a new panel for a new flow

Again:

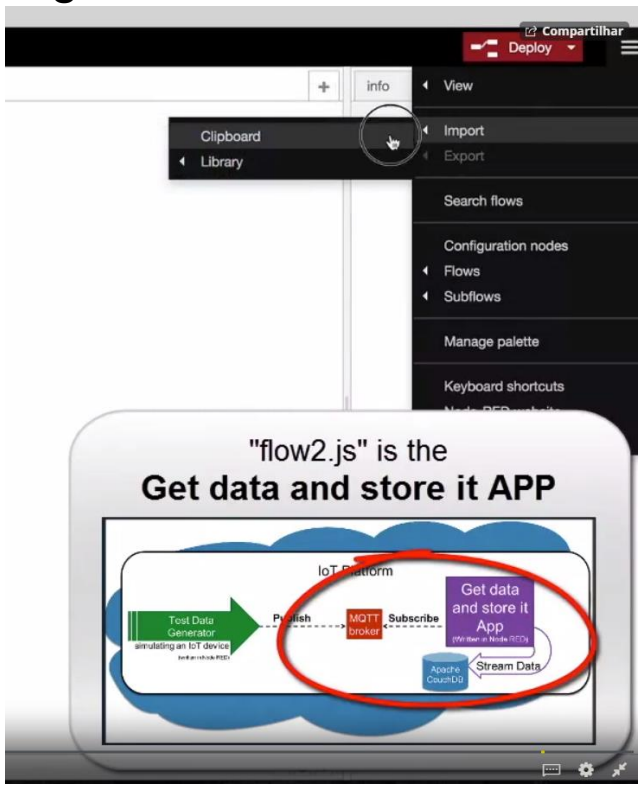

IBM IOT AND WASHING (FLOW2)

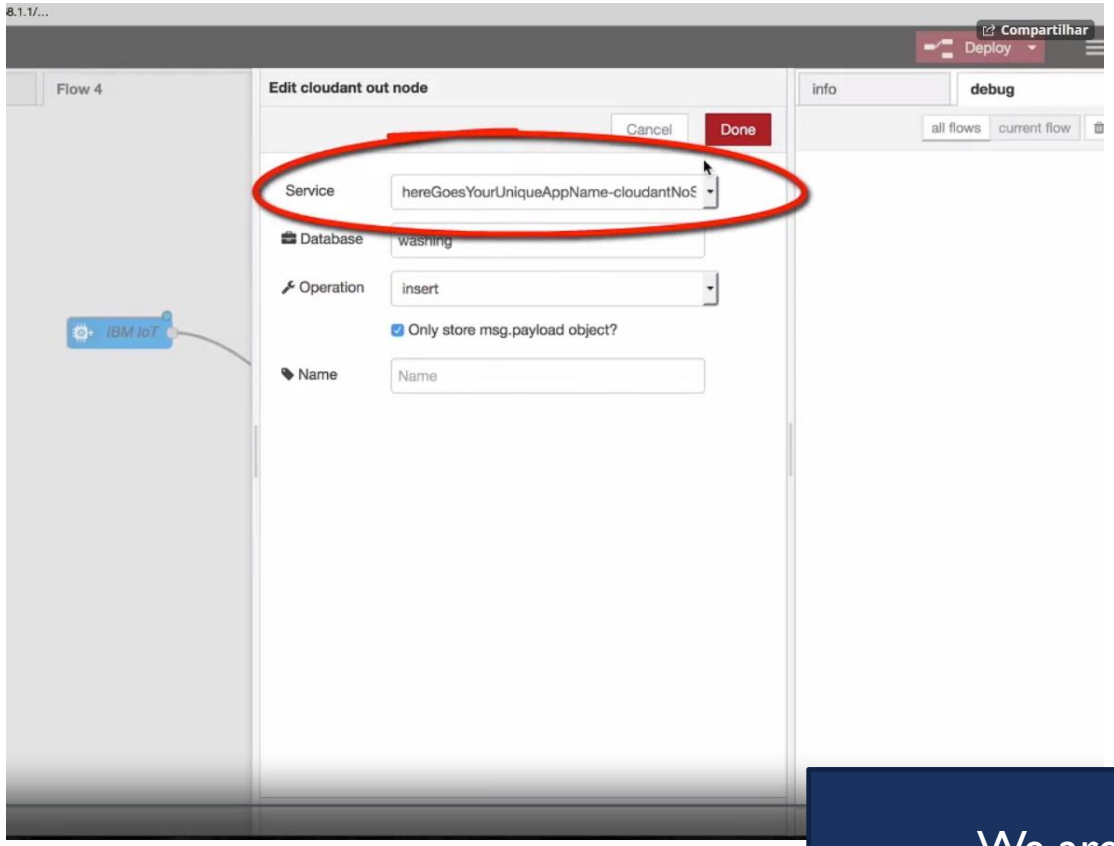

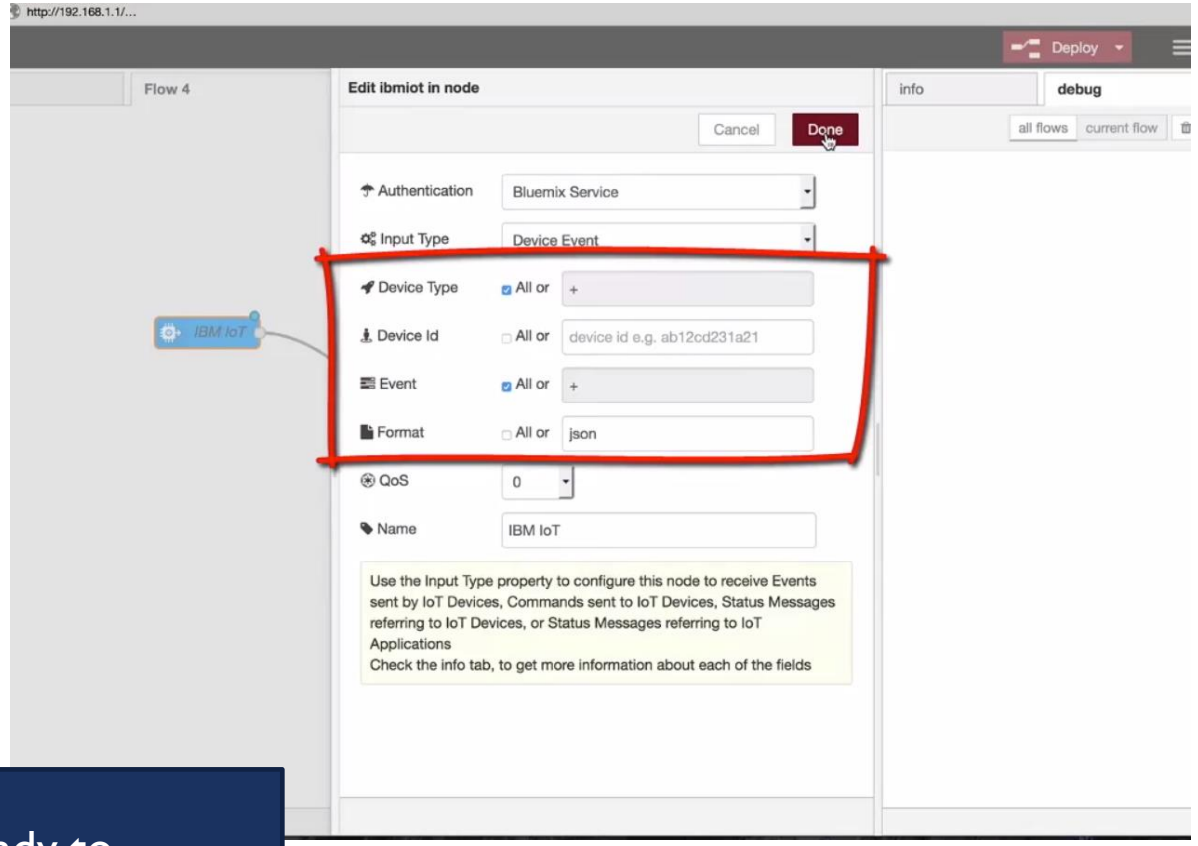

We are ready to
Generate test data!

CLICK ON MSG.PAYLOAD

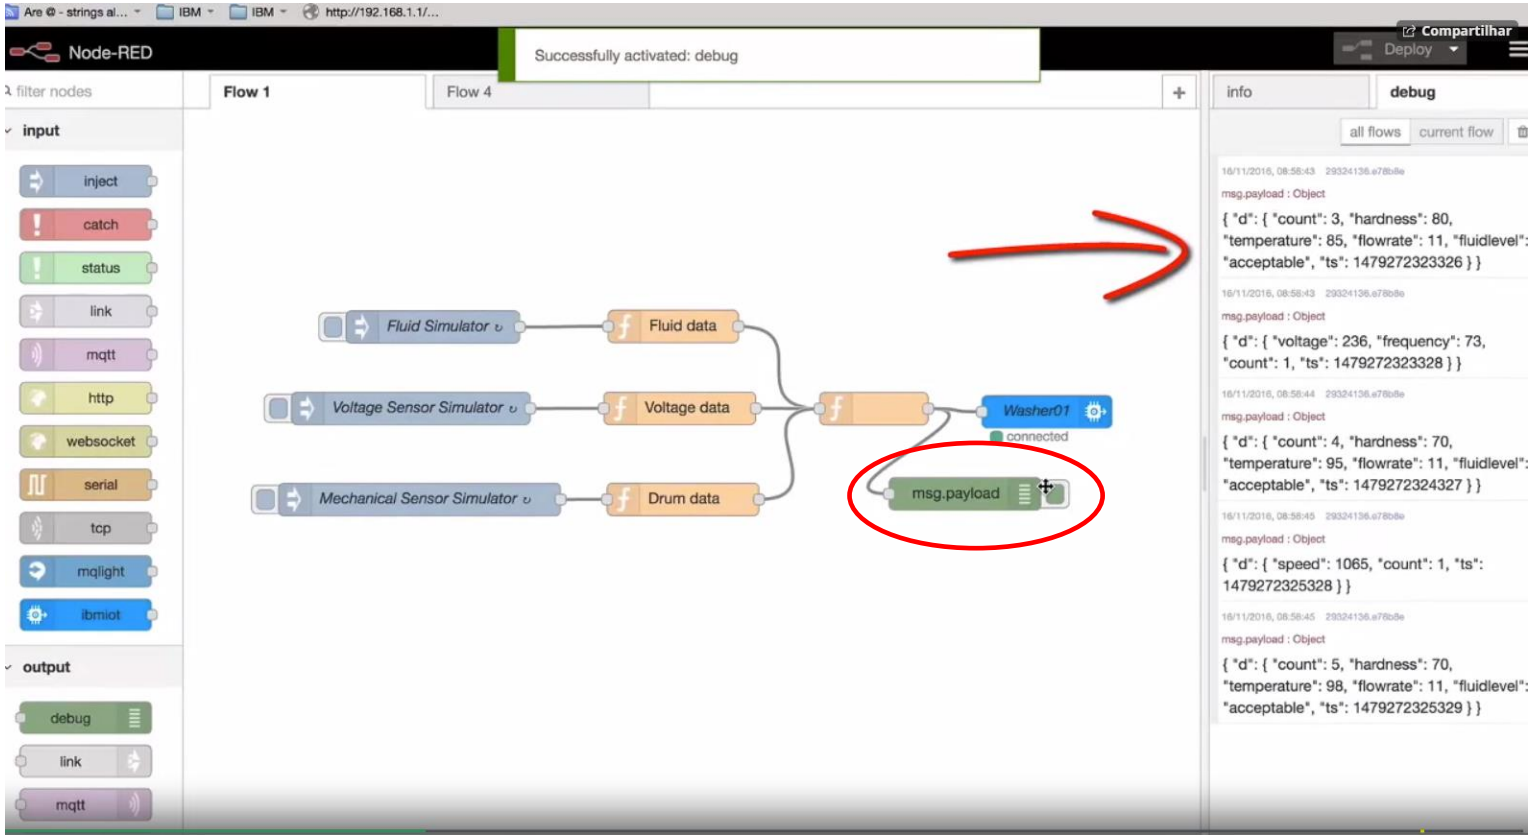

Once we activate the debug node, we will see test data generated by the test data generator.
We'll have a look at this later.
But basically, it is important to notice that this particular flow normally doesn't run in the cloud, but on an IoT device or gateway reading raw sensor data
directly from the built-in sensors and transmitting those directly to the cloud.
This is hypothetical sensor data coming from a washing machine.

Let's deactivate the debug node and add one on the cloud side.

COPY FOR THE FLOW 1 TO THE FLOW 2 (MSG.PAYLOAD)

Click here to start:

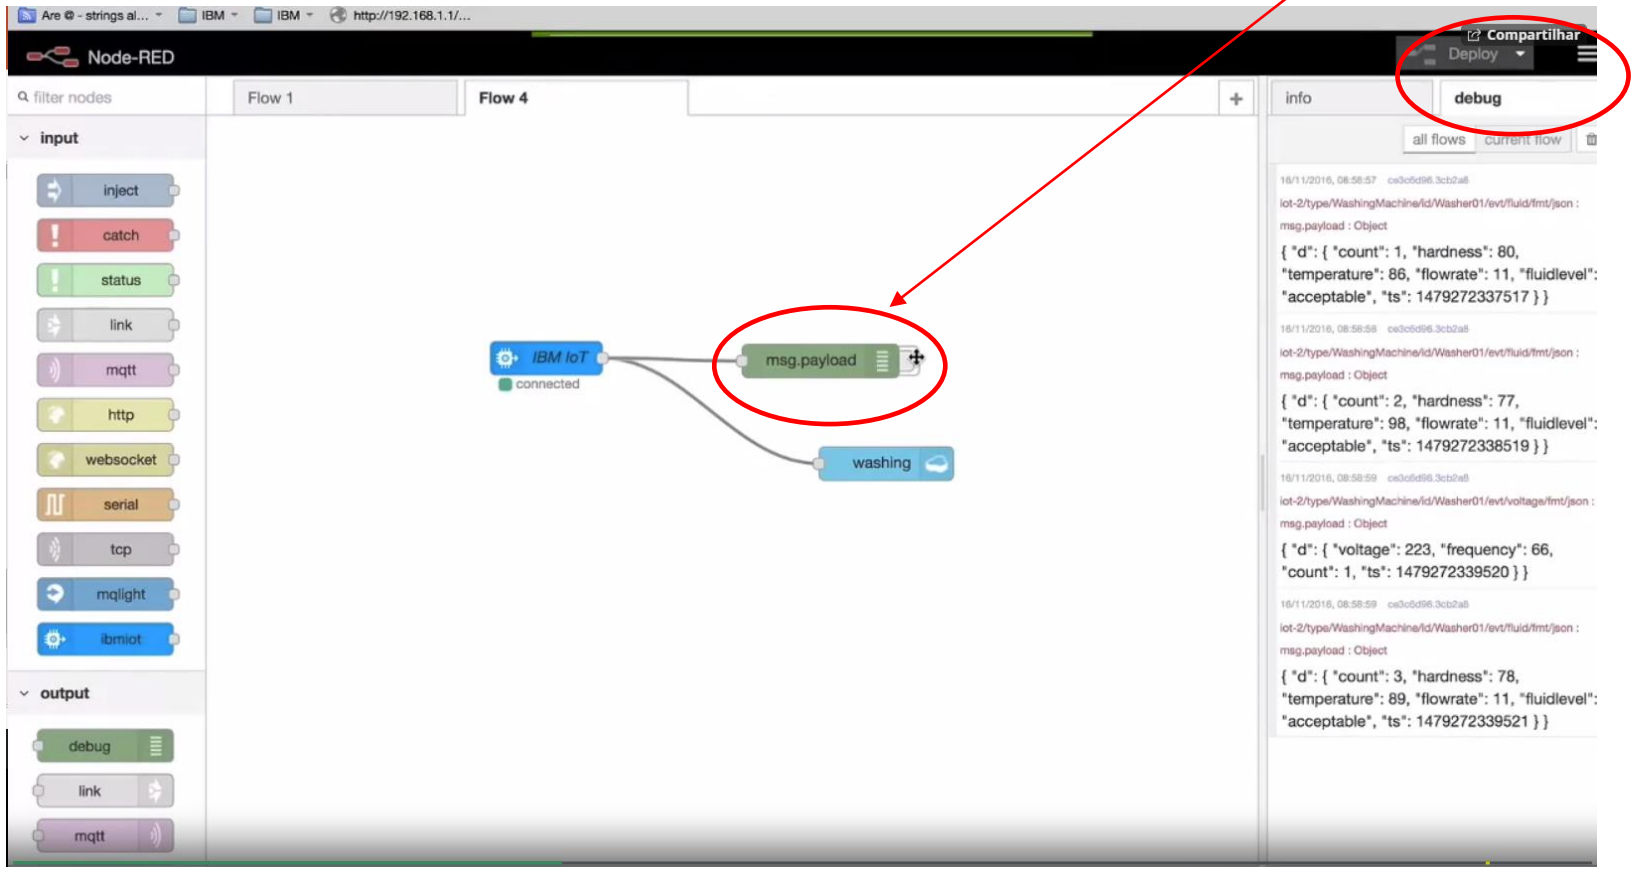

Now we can see the very same data, but now we are officially on the cloud and not pretending we are on an IoT device anymore.
We are currently streaming those data into Cloudant at the speed of more than one record per second.

SO NOW LET'S HAVE A BRIEF LOOK AT THE TEST DATA GENERATOR (FLOW1):

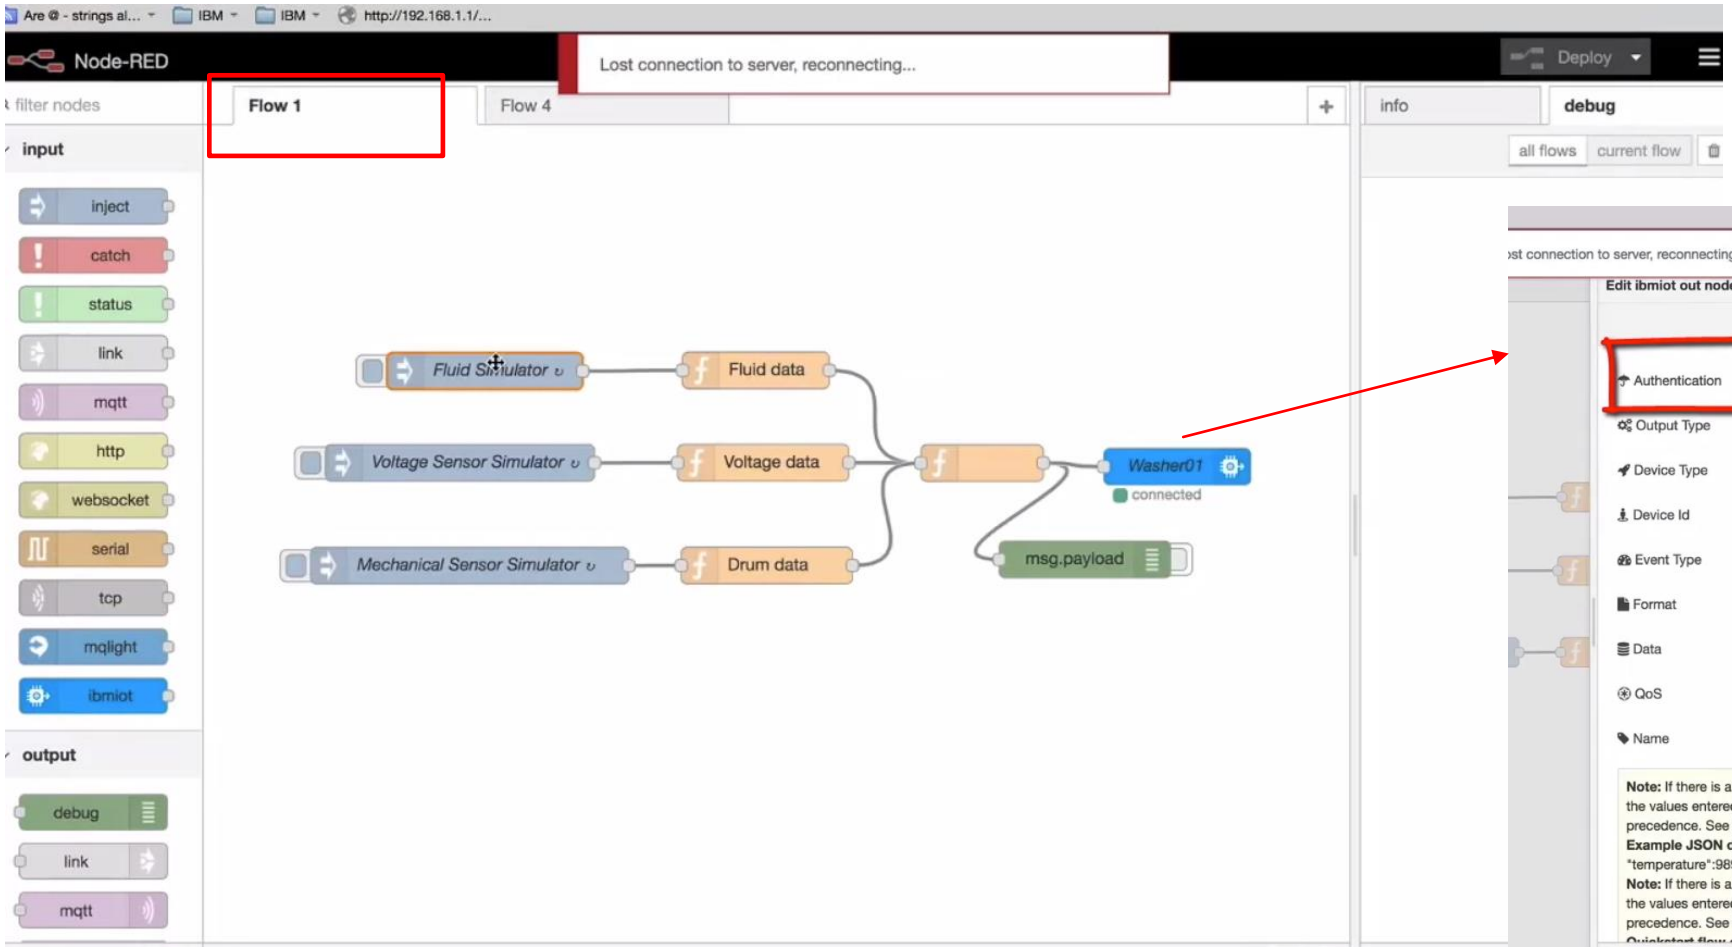

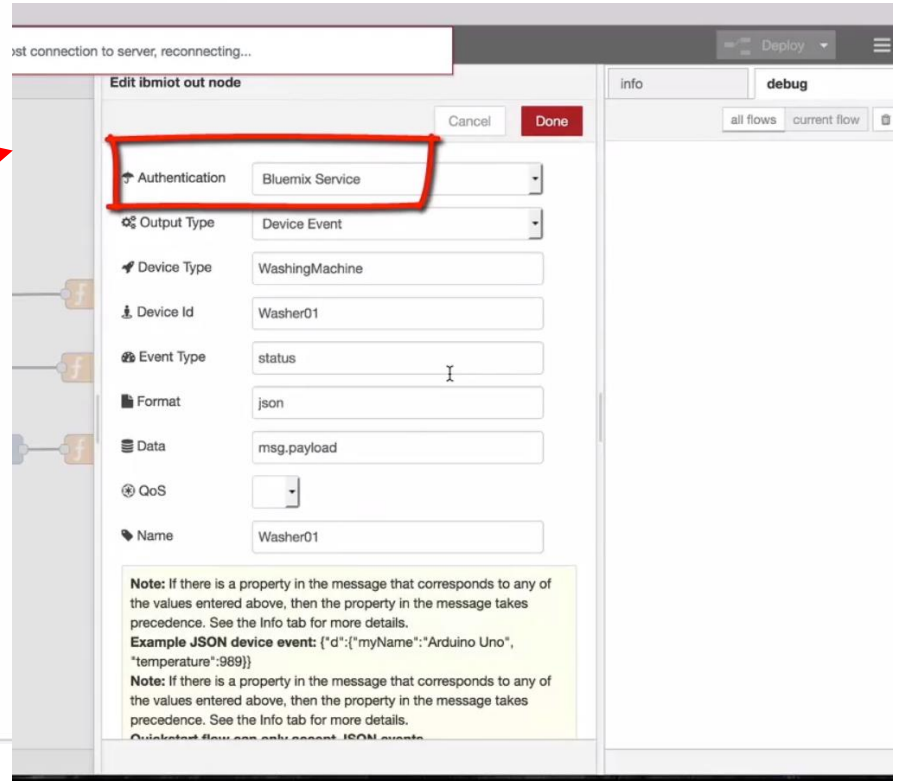

LET'S HAVE A LOOK AT THE CLOUDANT USER INTERFACE TO SEE WHETHER THE DATA WE ARE STORING ACTUALLY ARRIVES IN THE DATABASE.

- So we open the Bluemix user interface and
- scroll down to the Cloudant service which is bound to our Node-RED application.
- You can easily make autocorrect one since by default it has the same name as the URL of your application.

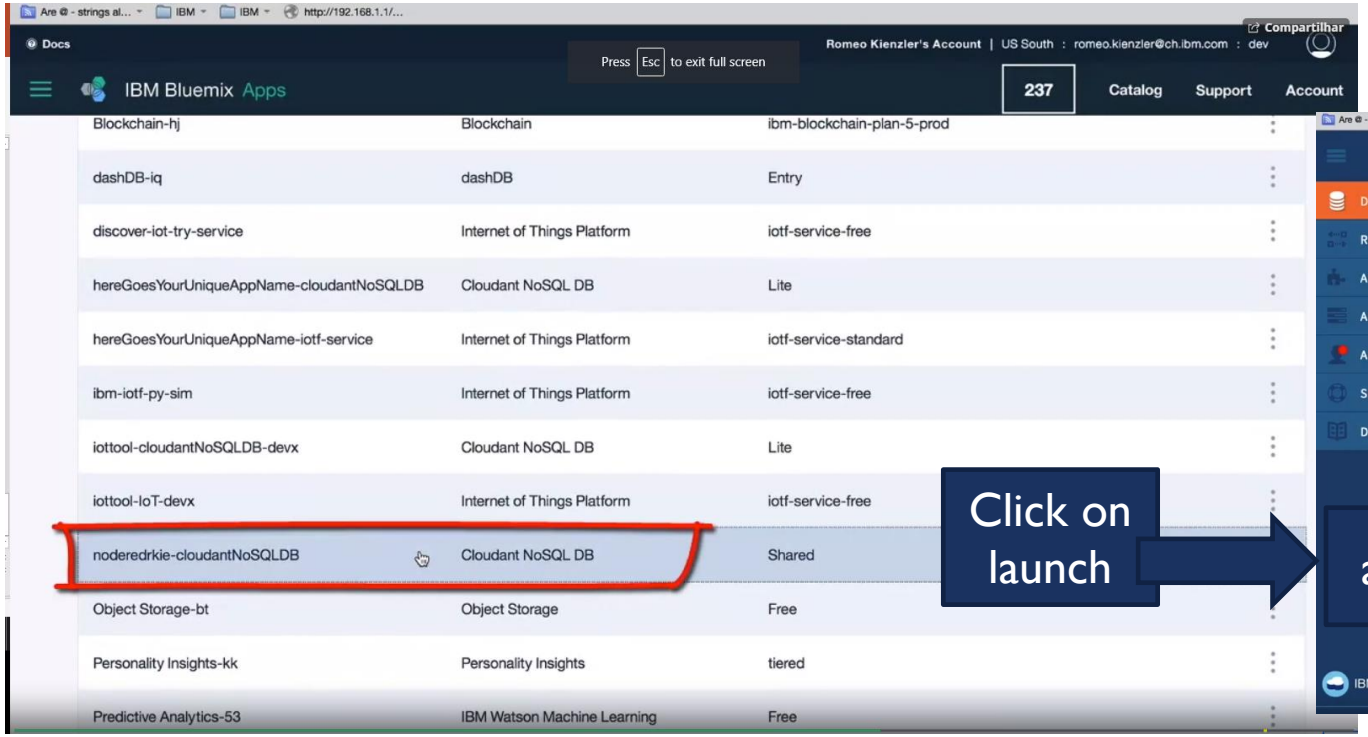

ACCESS CLOUDANT DATA ON APACHE SPARK

- In order to access Cloudant data on Apache Spark, we need to obtain the Cloudant credentials.
- So again, we enter the IBM Bluemix dashboard, click on our Node-RED application.

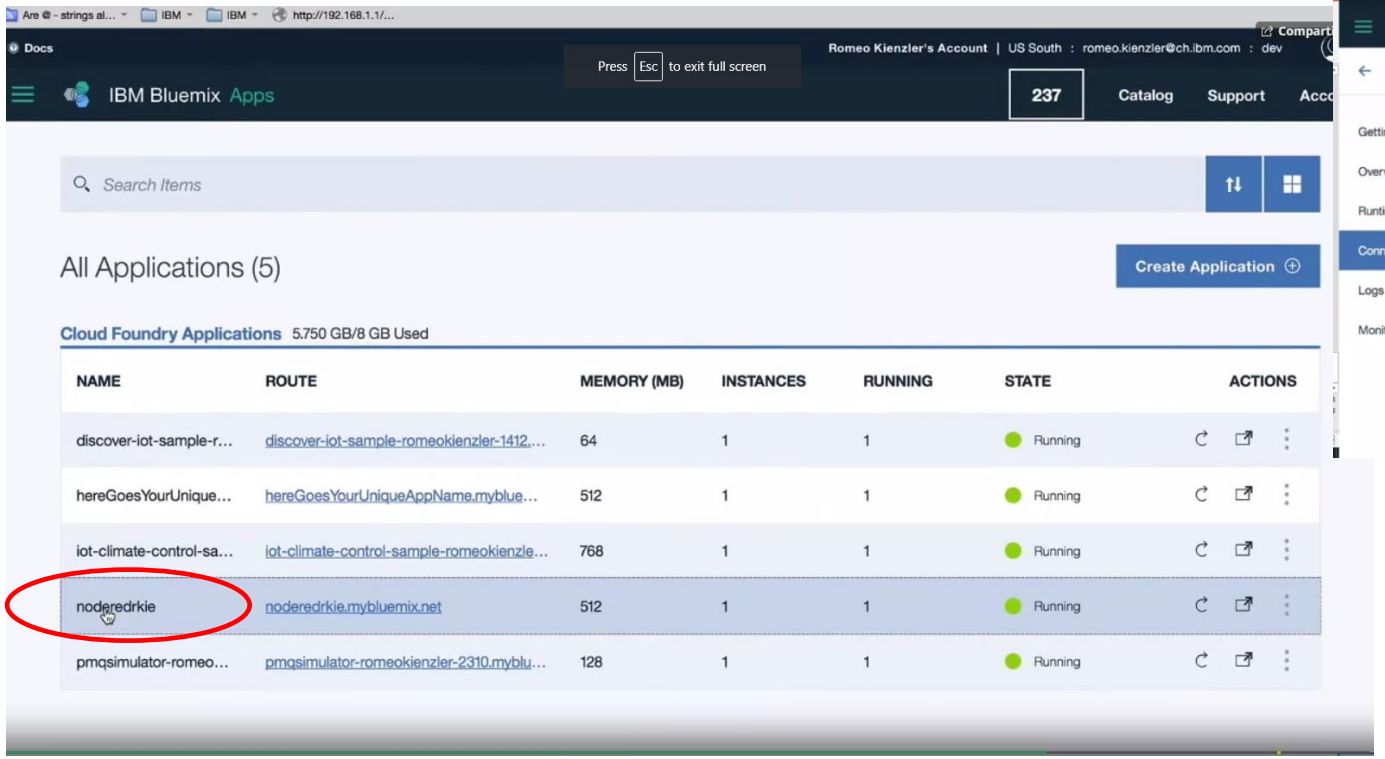

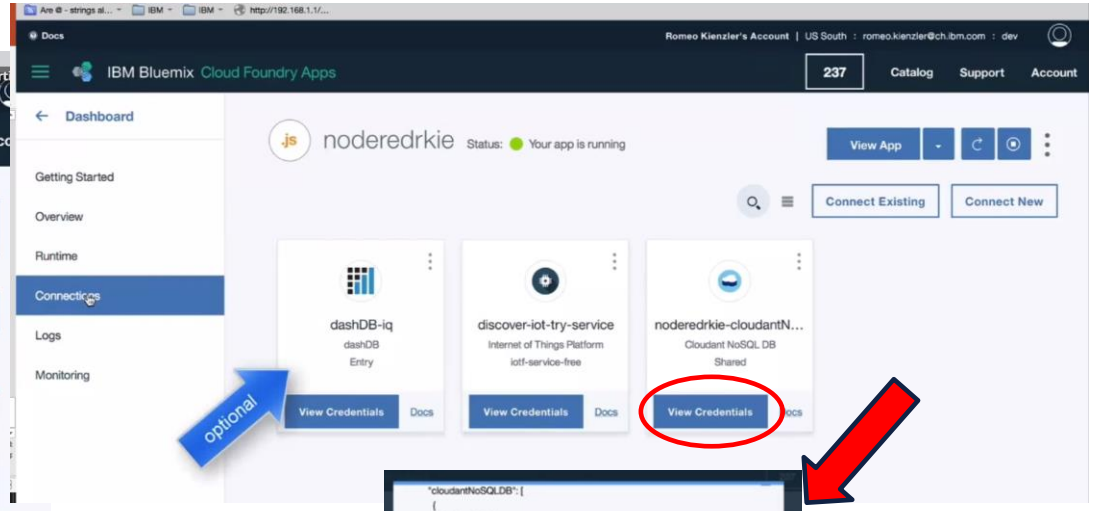

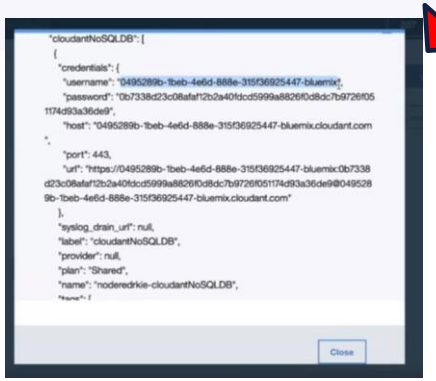

SO NOW IF WE HAVE ALL INFORMATION TO ACCESS THE DATABASE FROM APACHE SPARK, SO LET'S HAVE A LOOK HOW THIS WORKS.

- There are basically four parameters you have to specify when accessing Cloudant from Apache Spark using the Cloudant connector:

Flow 1" contains a NodeRED application to simulate IoT sensor data. This flow is running in the cloud. In a real scenario, where would this flow normally run?
- On a IoT Gateway and
- On a IoT Device!

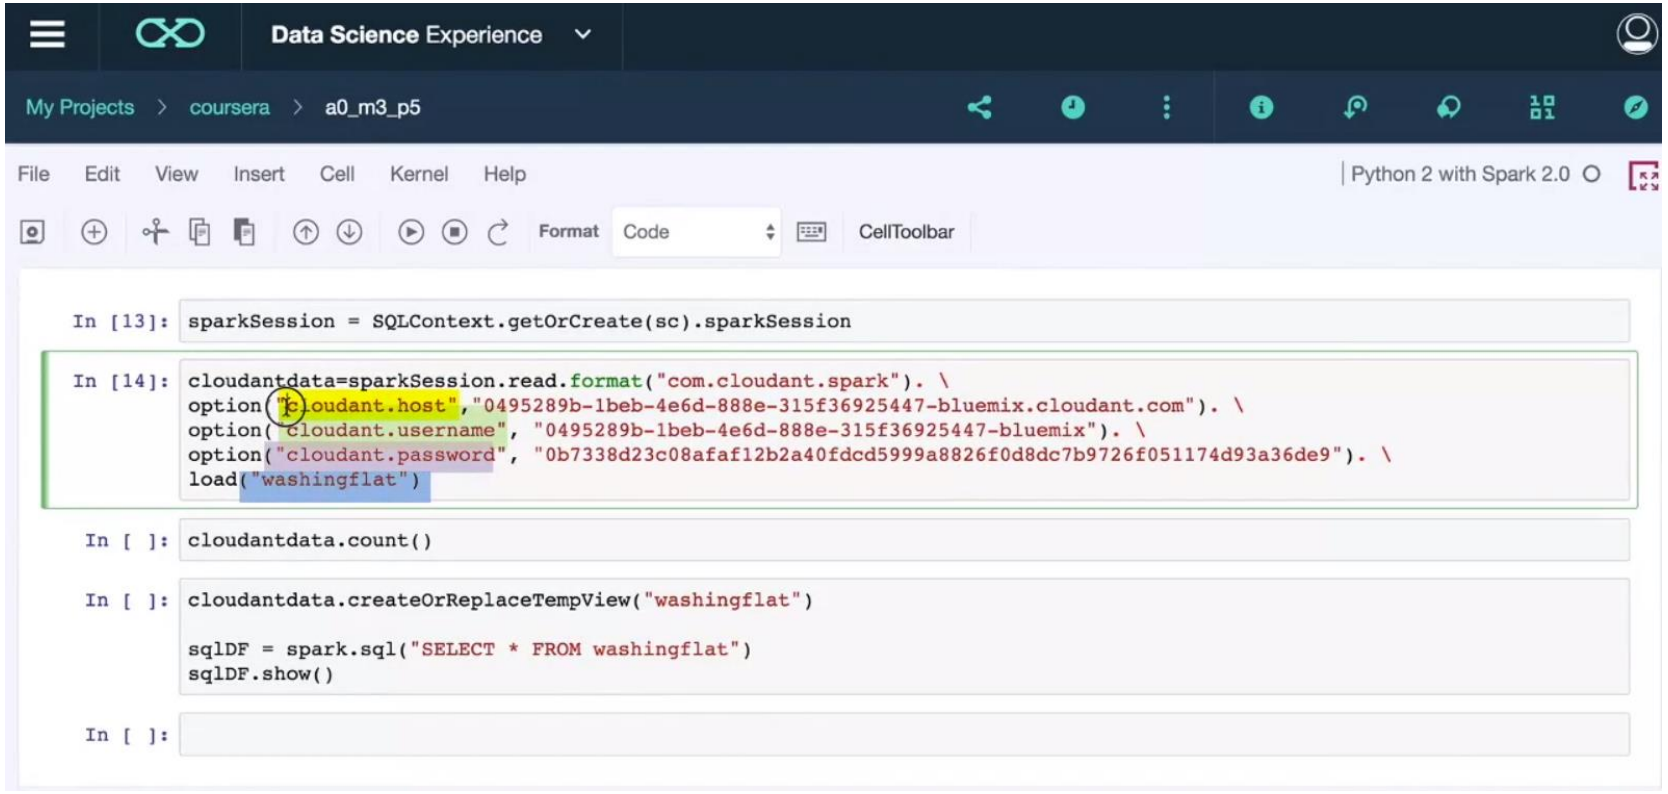

SUMMARY

- IoT data comes from IoT devices
- NodeRED can run on IoT devices
- Easy to simulate test data and run your solution – "the how"
- Straightforward to stream data to Cloudant
- ApacheSpark cloudant connector supports dataframe from Apache Spark